

Discussion Session 8:

AVL Trees

CS 225: Data Structures
& Software Principles

Agenda


- **Unbalanced Trees**
- AVL Trees
- Balancing at each Node
- Rotations
- Operations on AVL Trees
 - Insert
 - Remove
- Implementation of AVL Trees

By the end of this class, you

- Need to
 - Understand why balanced trees are desirable
 - Be able to calculate balance at each node in an AVL tree
 - Understand the 4 different types of rotations and their effects
 - Be able to restructure a tree after insertion/deletion
- Ought to be able to implement an AVL Tree

Motivation: BST

- Average case is $O(\log n)$





- But worst case search is $O(n)$!

Rotation helps!



Single Left and Single Right Rotations




Balanced Trees

- By mechanisms similar to the rotations we've just seen, you can always prevent a tree from getting too unbalanced.
- Different kinds of balanced Trees
 - AVL Trees
 - Red-Black Trees
 - *Splay Trees*

Agenda

- Unbalanced Trees
- **AVL Trees**
- Balancing at each Node
- Rotations
- Operations on AVL Trees
 - Insert
 - Remove
- Implementation of AVL Trees

AVL trees



- Named after creators: Adelson-Velskii and Landis
- Additional constraint to BST: balance
 - At each node, subtree heights differ by no more than 1.
- Every AVL Tree with n nodes has height less than $1.44 \log n$ (i.e. $O(\log n)$)
- Worst case search, insert and remove is $O(\log n)$


Agenda

- Unbalanced Trees
- AVL Trees
- **Balancing at each Node**
- Rotations
- Operations on AVL Trees
 - Insert
 - Remove
- Implementation of AVL Trees

AVL Tree Node Balance

- $\text{LeftHeight}(T) = 0$ if $\text{LeftChild}(T) == \text{NULL}$
 $1 + \text{Height}(T \rightarrow \text{Left})$ otherwise
- $\text{RightHeight}(T) = 0$ if $\text{RightChild}(T) == \text{NULL}$
 $1 + \text{Height}(T \rightarrow \text{Right})$ otherwise
- $\text{Height}(T) = \max \{ \text{LeftHeight}, \text{RightHeight} \}$
- $\text{Balance}(T) = \text{RightHeight}(T) - \text{LeftHeight}(T)$
- A node is only allowed to have a balance of:
-1, 0, +1

Example AVL Trees



■ 1

■ Height

■ Balance

■ 4

■ Height

■ Balance

■ 3

■ Height


■ Balance

■ 2

■ Height

■ Balance

Example AVL Trees



■ 1

■ Height 0

■ Balance 0

■ 8

■ Height 0

■ Balance 0

■ 7

■ Height 1


■ Balance 1

■ 2

■ Height 2

■ Balance 1

More Example Trees



■ $1 = ?, ?$

■ $6 = ?, ?$


■ $9 = ?, ?$

■ $8 = ?, ?$

■ $7 = ?, ?$

■ $2 = ?, ?$

More Example Trees



■ $1 = 0, 0$

■ $6 = 0, 0$


■ $9 = 0, 0$

■ $8 = 1, 1$

■ $7 = 2, 1$

■ $2 = 3, \underline{2}$

More Example Trees



■ $2 = ?, ?$

■ $1 = ?, ?$

■ $4 = ?, ?$

■ $3 = ?, ?$

■ $5 = ?, ?$

■ $7 = ?, ?$

■ $9 = ?, ?$

■ $13 = ?, ?$


■ $12 = ?, ?$

■ $15 = ?, ?$

■ $8 = ?, ?$

■ $6 = ?, ?$

More Example Trees



■ $2 = 0, 0$

■ $1 = 1, 1$

■ $4 = 0, 0$

■ $3 = 2, -1$

■ $5 = 3, -3$

■ $7 = 0, 0$

■ $9 = 0, 0$

■ $13 = 0, 0$

■ $12 = 1, 0$

■ $15 = 2, -2$

■ $8 = 3, 2$

■ $6 = 4, 0$

AVL Trees

- Binary Search Tree
- Operations similar to those in BSTs except
 - Remove & Insert
- Remove and Insert must make sure the tree stays balanced!
- Require re-balancing (Rotation) in some cases

Agenda

- Unbalanced Trees
- AVL Trees
- Balancing at each Node
- **Rotations**
- Operations on AVL Trees
 - Insert
 - Remove
- Implementation of AVL Trees


Rotations

- After running a BST insert/remove, must rebalance using rotations (rotate when the balance becomes -2 or $+2$).
- Must know:
 - *HOW?* how to apply a rotation
 - *WHICH?* which type of rotation to use
 - *WHERE?* where to apply the rotation


How?

- There are 4 different rotations:
 - single right
 - single left
 - double rotation LR
 - double rotation RL
- If you know single right, you can derive the other three!


Single Right Rotation



Single Left Rotation




Double Rotation (LR)



This is an LR Double rotation : requires a single Left followed by a single Right rotation


Double Rotation (RL)



This is an RL Double rotation : requires a single Right followed by a single Left rotation

Which?

- Find the new balance for the critical node and its child on the path to the inserted/deleted node.
- If these balance values are of the same sign, then you need only a single rotation. Otherwise, you need to do a double rotation.
- In case of a single rotation, additionally, if the sign on the critical node is +, do a left rotation and if it is -, do a right rotation.
- In case of a double rotation, additionally, if the sign on the critical node is +, do an RL and if it is -, do an LR.



Which?

- Type: Find the new balance for the critical node and child on the path to the inserted/deleted node.
 - If balance values are the *same sign*, then you need only a **single** rotation.
 - Else, you need to do a **double** rotation.
- Direction of Single Rotation:
 - if the sign on the critical node is **+**, do a **left** rotation
 - else if it is **-**, do a **right** rotation
- Direction of Double Rotation:
 - if the sign on the critical node is **+**, do an RL (double left)
 - else if it is **-**, do an LR (double right)

Where?

- When you insert or remove a node, height *may* change and balance *will* change.
- As you move back up the tree, must adjust heights and calculate a new balance.
- If new balance is +2 or -2, *this* is the node around which to perform a rotation.

Rotation Summary

- Memorize single right rotation
- Single left is mirror image of single right
- Double rotations are a combination of two single rotations:
 - 1st on child in first direction
 - 2nd on **critical node** in second direction

Rotation Effects

- Reduce balance of critical node from +2 or -2 to 0
- Reduce height of critical node by 1

Agenda

- Unbalanced Trees
- AVL Trees
- Balancing at each Node
- Rotations
- **Operations on AVL Trees**
 - Insert
 - Remove
- Implementation of AVL Trees

Insert

- Perform BST Insert
- New node is a leaf and has height 0, balance 0
- Work way up tree and compute new balance (starting with new node's parent)
 - if balance goes from ± 1 to 0, stop.
 - if balance goes from 0 to ± 1 , then height has also changed. Proceed up.
 - if balance goes from +1 to +2 or -1 to -2, do a rotation on this node! ... then stop.

Remove

- Perform BST remove (in-order successor or predecessor replaces removed node).
- Compute new balance of parent:
 - if balance goes from 0 to ± 1 ...done
 - if balance goes from ± 1 to 0, height has also changed. Proceed up.
 - If balance goes from +1 to +2 or -1 to -2, do a rotation. Height has also changed, proceed up.

Operation Summary

- Insert has at most one single or one double rotation.
- Remove can have a rotation at every level.

Agenda


- Unbalanced Trees
- AVL Trees
- Balancing at each Node
- Rotations
- Operations on AVL Trees
 - Insert
 - Remove
- **Implementation of AVL Trees**

Source


- AVL Tree Code found at:
~cs225/src/library/08-avl/
 - Designed in the spirit of C++ STL interfaces
 - typedefs, 4 types of iterators, functions to return iterators, regular operations
 - Stores key-element pairs
(template <class Ktype, class Etype>
 - Code that follows may not contain all the gory details!
- The (not very good) applet can be found at
<http://www.binarytreesome.com/applets/latest/applet.html>

class avl_tree Private Data

```
class avl_tree_node {  
    public:  
        avl_tree_node() : element(), left(NULL), right(NULL), parent(NULL),  
                           height(0) {}  
        avl_tree_node(pair<Ktype, Etype> elmt,  
                       avl_tree_node* leftPtr = NULL, avl_tree_node* rightPtr = NULL,  
                       avl_tree_node* parentPtr = NULL, int hgt = 0) ;  
  
        pair<Ktype, Etype> element;           // value element of node  
        avl_tree_node* left;                  // pointer to left subtree  
        avl_tree_node* right;                 // pointer to right subtree  
        avl_tree_node* parent;                // pointer to parent  
        int height;                           // height of node  
};  
avl_tree_node* headerNode;                   // pointer to header node [ end() ]  
int treeSize;                                // number of nodes in tree
```



headerNode




- `headerNode->parent` = root node [`root()`]
- `headerNode->left` = leftmost key in tree [`leftmost()`]
- `headerNode->right` = rightmost key in tree [`rightmost()`]
- Useful for working with iterators

single_left_rotation()

// Performs a rotation between a node (k2) and its right child, updating the
// heights of both. Can be called only if k2 has a right child.


```
void single_left_rotation( avl_tree_node* & k2 ) {  
    avl_tree_node *k1 = k2->right;  
    avl_tree_node *parentOfk2 = k2->parent;  
  
    // rotate k1 up to "root" and k2 down to left  
    k2->right = k1->left;  
    if (k2->right != NULL)    k2->right->parent = k2;  
    k1->left = k2;  
    k2->parent = k1;  
  
    k1->parent = parentOfk2;  
  
    calculate_height(k2);  
    calculate_height(k1);  
    // reset the root  
    k2 = k1;  
}
```



double_left_rotation()

// Performs a right-left rotation, also known as a double left rotation. Can
// only be called if k3 has a right child and k3's right child has a left child.

```
void double_left_rotation( avl_tree_node* & k3 ) {  
    single_right_rotation(k3->right);  
    single_left_rotation(k3);  
}
```



Public insert() Interface

// - return value : a pair that contains an iterator to the inserted
// element and a boolean integer telling us if the insertion was
// performed (1) or if we instead were trying to insert a duplicate
// key and thus stopped (0).

```
pair<iterator, int> insert(const pair<Ktype, Etype>& insElem) {  
    if (count(key(insElem)) == 1)      // this key already appears  
        // note the constructor call  
        return pair<iterator, int>( find(key(insElem)), 0);  
    else  
        return pair<iterator, int>(  
            insert(insElem, root(), headerNode), 1);  
}
```

private insert()

```
iterator insert(const pair& insElem, avl_tree_node*& TN, avl_tree_node* parentOfTN) {  
    // we've found the spot to insert at  
    if (TN == NULL) {  
        // Create and insert the node with value insElem  
        TN = new avl_tree_node();  
        TN->element = insElem;  
        TN->height = 0;  
        TN->parent = parentOfTN;  
  
        // new code  
        if ( (treeSize==0) || (key(insElem) < key(headerNode->left->element)) )  
            headerNode->left = TN;  
        if ( (treeSize==0) || (key(insElem) > key(headerNode->right->element)) )  
            headerNode->right = TN;  
  
        treeSize++;  
        return iterator(TN);  
    }  
}
```

...private **insert** continued...

```
else if (key(insElem) < key(TN->element) ) {  
    // insert in the left subtree  
    iterator tempIt = insert(insElem, TN->left, TN);  
    // check balance condition  
    if ( node_height(TN->left) - node_height(TN->right) == 2) {  
        // if unbalanced, determine type of rotation  
        if (key(insElem) < key(TN->left->element) ) {  
            // if inserted in left-most subtree, do single  
            single_right_rotation(TN);  
        } else {  
            // if inserted in interior subtree, do double  
            double_right_rotation(TN);  
        }  
    } else {  
        // if balanced, recalculate height (due to an insertion further down the tree)  
        calculate_height(TN);  
    }  
    return tempIt;  
}
```

...private insert() finished.

```
else if (key(insElem) > key(TN->element) ) {  
    // insert into right subtree...  
  
    // ...same as previous slide, but anywhere there's a "left" replace with  
    // "right" and vice versa.  
}  
// else insElem is in the tree already--do nothing  
else  
    return end( );
```

private remove() components

```
void remove(pair & remElem, avl_tree_node* & TN, avl_tree_node* parentOfTN) {  
    avl_tree_node *tmp_cell, *tmpCellParent, *newCell;    // used in 2 child remove case  
    if (TN == NULL)  
        return;  
    else if ( key(remElem) < key(TN->element) ) {  
        // remove from the left subtree  
        remove(remElem, TN->left, TN);  
        // check balance condition, perform left rotation if necessary, update heights  
        // ...  
    }  
    else if (key(remElem) > key(TN->element)) {  
        // remove from the right subtree  
        remove(remElem, TN->right, TN);  
        // check balance condition, perform right rotation if necessary, update heights  
        // ...  
    }  
    else { // remElem == TN->element  
        // 2 child case, 1 child case, 0 child case, fix headerNode if necessary  
        // see next slides...
```

...private remove() 2-child case...

```
// 2 child case -- replace TN (node we want to delete) with its inorder successor node
```

```
if ((TN->left != NULL) && (TN->right != NULL) ) {
```

```
    // find inorder successor
```

```
    tmp_cell = TN->right; // we just checked, TN->right isn't NULL
```

```
    while (tmp_cell->left != NULL)
```

```
        tmp_cell = tmp_cell->left;
```

```
    // Put IOS node in TN's place, and put a new node in IOS's place. We do this to:
```

```
    // 1) make sure any iterators pointing to IOS are still valid
```

```
    // 2) allow recursive calls to remove() to delete the new IOS and balance the tree
```

```
    tmpCellParent = tmp_cell->parent;
```

```
    newCell = new avl_tree_node();
```

```
    // make newCell a copy of tmp_cell
```

```
    newCell->right = tmp_cell->right;
```

```
    newCell->left = tmp_cell->left;
```

```
    newCell->element = tmp_cell->element;
```

```
    newCell->height = tmp_cell->height;
```

```
    newCell->parent = tmp_cell->parent;
```

...remove() 2-child case finished...

// if inorder successor is left child of its parent

if (tmpCellParent->left == tmp_cell) tmpCellParent->left = newCell;

// else if inorder successor is right child of its parent

else if (tmpCellParent->right == tmp_cell) tmpCellParent->right = newCell;

// now, newCell has completely taken the place of tmp_cell.

// make tmp_cell a copy of TN, except TN->element – we want to remove that!

tmp_cell->right = TN->right;

tmp_cell->parent = TN->parent;

tmp_cell->left = TN->left;

tmp_cell->height = TN->height;

delete TN; // okay, since tmp_cell holds all data

TN = tmp_cell;

TN->left->parent = TN;

TN->right->parent = TN;

// now, finally, **delete the just added new tmp_cell**

remove(TN->element, TN->right, TN);

// **check balance condition**, perform rotation if necessary, update heights ...

} // end 2 child case

...remove() 1-child & 0-child cases

```
else {  
    // X has 0 children or 1 child  
    tmp_cell = TN;  
    if ((TN->left == NULL) && (TN->right == NULL)) {    // no children  
        TN = NULL;  
    }  
    else if ((TN->left == NULL) && (TN->right != NULL)) { // no left child  
        TN = TN->right;  
        TN->parent = parentOfTN;  
    }  
    else if ((TN->left != NULL) && (TN->right == NULL)) { // no right child  
        TN = TN->left;  
        TN->parent = parentOfTN;  
    }  
    delete tmp_cell;  
    // reassign leftmost/rightmost in headerNode if need be ...  
    treeSize--;  
} // end 1-child & 0-child case  
} // remElem == TN->element  
} // end remove()
```


Summary

- Unbalanced Trees
- AVL Trees
- Balancing at each Node
- Rotations
- Operations on AVL Trees
 - Insert
 - Remove
- Implementation of AVL Trees